

- ✓ 7. Square Trick
- ✓ 8. Gradient Descent
- ✓ 9. Mean Absolute Error
- ✓ 10. Mean Squared Error
- ✓ 11. Minimizing Error Functions
- ✓ 12. Mean vs Total Error
- ✓ 13. Mini-batch Gradient Descent
- ✓ 14. Absolute Error vs Squared Error
- ✓ 15. Linear Regression in scikit-learn
- ✓ 16. Higher Dimensions
- ✓ 17. Multiple Linear Regression
- 18. Closed Form Solution
- 19. (Optional) Closed form Solution ...
- 20. Linear Regression Warnings



Multiple Linear Regression

Multiple Linear Regression

In the last section, you saw how we can predict life expectancy using BMI. Here, BMI was the **predictor**, also known as an independent variable. A predictor is a variable you're looking at in order to make predictions about other variables, while the values you are trying to predict are known as dependent variables. In this case, life expectancy was the dependent variable.

Now, let's say we get new data on each person's heart rate as well. Can we create a prediction of life expectancy using both BMI and heart rate?

Absolutely! As we saw in the previous video, we can do that using multiple linear regression.

If the outcome you want to predict depends on more than one variable, you can make a more complicated model that takes this into account. As long as they're relevant to the situation, using more independent/predictor variables can help you get a better prediction.

When there's just one predictor, the linear regression model is a line, but as you add more predictor variables, you're adding more dimensions to the picture.

When you have one predictor variable, the equation of the line is

$$y = mx + b$$



Multiple Linear Regression



Linear regression with one predictor variable


Adding a predictor variable to go to two predictor variables means that the predicting equation is:

$$y = m_1x_1 + m_2x_2 + b$$

To represent this graphically, we'll need a three-dimensional plot, with the linear regression model represented as a plane:



Multiple Linear Regression



Linear regression with two predictor variables

You can use more than two predictor variables - in fact, you should use as many as is useful! If you use n predictor variables, then the model can be represented by the equation

$$y = m_1x_1 + m_2x_2 + m_3x_3 + \dots + m_nx_n + b$$

As you make a model with more predictor variables, it becomes harder to visualise, but luckily, everything else about linear regression stays the same. We can still fit models and make predictions in exactly the same way - time to try it!



Multiple Linear Regression

consists of 13 features of 506 houses and the median home value in \$1000's. You'll fit a model on the 13 features to predict the value of the houses.

You'll need to complete each of the following steps:

1. Build a linear regression model

- Create a regression model using scikit-learn's **LinearRegression** and assign it to **model**.
- Fit the model to the data.

2. Predict using the model

- Predict the value of **sample_house**.

quiz.py

solution.py

```
1 from sklearn.linear_model import LinearRegression
2 from sklearn.datasets import load_boston
3
4 # Load the data from the boston house-prices dataset
5 boston_data = load_boston()
6 x = boston_data['data']
7 y = boston_data['target']
8
9 # Make and fit the linear regression model
10 # TODO: Fit the model and assign it to the model variable
11 model = LinearRegression()
12 model.fit(x,y)
13
14 # Make a prediction using the model
15 sample_house = [[2.29690000e-01, 0.00000000e+00, 1.05900000e+01, 0.000
16                  6.32600000e+00, 5.25000000e+01, 4.35490000e+00, 4.000
17                  1.86000000e+01, 3.94870000e+02, 1.09700000e+01]]
```



Multiple Linear Regression

23.6842056923 is the correct prediction!

RESET QUIZ

TEST RUN

SUBMIT ANSWER

NEXT