

Linear Regression

- [Introduction](#)
- [Simple regression](#)
 - [Making predictions](#)
 - [Cost function](#)
 - [Gradient descent](#)
 - [Training](#)
 - [Model evaluation](#)
 - [Summary](#)
- [Multivariable regression](#)
 - [Growing complexity](#)
 - [Normalization](#)
 - [Making predictions](#)
 - [Initialize weights](#)
 - [Cost function](#)
 - [Gradient descent](#)
 - [Simplifying with matrices](#)
 - [Bias term](#)
 - [Model evaluation](#)

Introduction

Linear Regression is a supervised machine learning algorithm where the predicted output is continuous and has a constant slope. It's used to predict values within a continuous range, (e.g. sales, price) rather than trying to classify them into categories (e.g. cat, dog). There are two main types:

Simple linear regression uses traditional slope-intercept form, where m and b are the variables our algorithm will try to “learn” to produce the most accurate predictions. x represents our input data and y represents our prediction.

$$y = mx + b$$

Multivariable regression

A more complex, multi-variable linear equation might look like this, where w represents the coefficients, or weights, our model will try to learn.

$$f(x, y, z) = w_1x + w_2y + w_3z$$

The variables x, y, z represent the attributes, or distinct pieces of information, we have about each observation. For sales predictions, these attributes might include a company’s advertising spend on radio, TV, and newspapers.

$$Sales = w_1Radio + w_2TV + w_3News$$

Simple regression

Let’s say we are given a [dataset](#) with the following columns (features): how much a company spends on Radio advertising each year and its annual Sales in terms of units sold. We are trying to develop an equation that will let us to predict units sold based on how much a company spends on radio advertising. The rows (observations) represent companies.

Company	Radio (\$)	Sales
Amazon	37.8	22.1
Google	39.3	10.4
Facebook	45.9	18.3

Apple	41.3	18.5
-------	------	------

Making predictions

Our prediction function outputs an estimate of sales given a company's radio advertising spend and our current values for *Weight* and *Bias*.

$$Sales = Weight \cdot Radio + Bias$$

Weight

the coefficient for the Radio independent variable. In machine learning we call coefficients *weights*.


Radio

the independent variable. In machine learning we call these variables *features*.

Bias

the intercept where our line intercepts the y-axis. In machine learning we can call intercepts *bias*. Bias offsets all predictions that we make.

Our algorithm will try to *learn* the correct values for Weight and Bias. By the end of our training, our equation will approximate the *line of best fit*.



Code

```
def predict_sales(radio, weight, bias):  
    return weight*radio + bias
```

Cost function

The prediction function is nice, but for our purposes we don't really need it. What we need is a [cost function](#) so we can start optimizing our weights.

Let's use [MSE \(L2\)](#) as our cost function. MSE measures the average squared difference between an observation's actual and predicted values. The output is a single number representing the cost, or score, associated with our current set of weights. Our goal is to minimize MSE to improve the accuracy of our model.

Given our simple linear equation $y = mx + b$, we can calculate MSE as:

$$MSE = \frac{1}{N} \sum_{i=1}^n (y_i - (mx_i + b))^2$$

! Note

- N is the total number of observations (data points)
- $\frac{1}{N} \sum_{i=1}^n$ is the mean
- y_i is the actual value of an observation and $mx_i + b$ is our prediction

Code

```
def cost_function(radius, sales, weight, bias):  
    companies = len(radius)  
    total_error = 0.0  
    for i in range(companies):  
        total_error += (sales[i] - (weight*radius[i] + bias))**2  
    return total_error / companies
```

Gradient descent

To minimize MSE we use [Gradient Descent](#) to calculate the gradient of our cost function. [TODO: Slightly longer explanation].

Math

There are two [parameters](#) (coefficients) in our cost function we can control: weight m and bias b . Since we need to consider the impact each one has on the final prediction, we use partial derivatives. To find the partial derivatives, we use the [Chain rule](#). We need the chain rule because

$(y - (mx + b))^2$ is really 2 nested functions: the inner function $y - (mx + b)$ and the outer function x^2 .

Returning to our cost function:

$$f(m, b) = \frac{1}{N} \sum_{i=1}^n (y_i - (mx_i + b))^2$$

We can calculate the gradient of this cost function as:

$$\begin{aligned} f'(m, b) &= \begin{bmatrix} \frac{df}{dm} \\ \frac{df}{db} \end{bmatrix} = \begin{bmatrix} \frac{1}{N} \sum -x_i \cdot 2(y_i - (mx_i + b)) \\ \frac{1}{N} \sum -1 \cdot 2(y_i - (mx_i + b)) \end{bmatrix} \\ &= \begin{bmatrix} \frac{1}{N} \sum -2x_i(y_i - (mx_i + b)) \\ \frac{1}{N} \sum -2(y_i - (mx_i + b)) \end{bmatrix} \end{aligned}$$

Code

To solve for the gradient, we iterate through our data points using our new weight and bias values and take the average of the partial derivatives. The resulting gradient tells us the slope of our cost function at our current position (i.e. weight and bias) and the direction we should update to reduce our cost function (we move in the direction opposite the gradient). The size of our update is controlled by the [learning rate](#).

```

def update_weights(radius, sales, weight, bias, learning_rate):
    weight_deriv = 0
    bias_deriv = 0
    companies = len(radius)

    for i in range(companies):
        # Calculate partial derivatives
        # -2x(y - (mx + b))
        weight_deriv += -2*radius[i] * (sales[i] - (weight*radius[i] + bias))

        # -2(y - (mx + b))
        bias_deriv += -2*(sales[i] - (weight*radius[i] + bias))

    # We subtract because the derivatives point in direction of steepest ascent
    weight -= (weight_deriv / companies) * learning_rate
    bias -= (bias_deriv / companies) * learning_rate

    return weight, bias

```

Training

Training a model is the process of iteratively improving your prediction equation by looping through the dataset multiple times, each time updating the weight and bias values in the direction indicated by the slope of the cost function (gradient). Training is complete when we reach an acceptable error threshold, or when subsequent training iterations fail to reduce our cost.

Before training we need to initialize our weights (set default values), set our [hyperparameters](#) (learning rate and number of iterations), and prepare to log our progress over each iteration.

Code

```

def train(radio, sales, weight, bias, learning_rate, iters):
    cost_history = []

    for i in range(iters):
        weight, bias = update_weights(radio, sales, weight, bias, learning_rate)

        #Calculate cost for auditing purposes
        cost = cost_function(radio, sales, weight, bias)
        cost_history.append(cost)

        # Log Progress
        if i % 10 == 0:
            print "iter={:d}    weight={:.2f}    bias={:.4f}    cost={:.2f}".format(i, weight, bias,
cost)

    return weight, bias, cost_history

```


Model evaluation

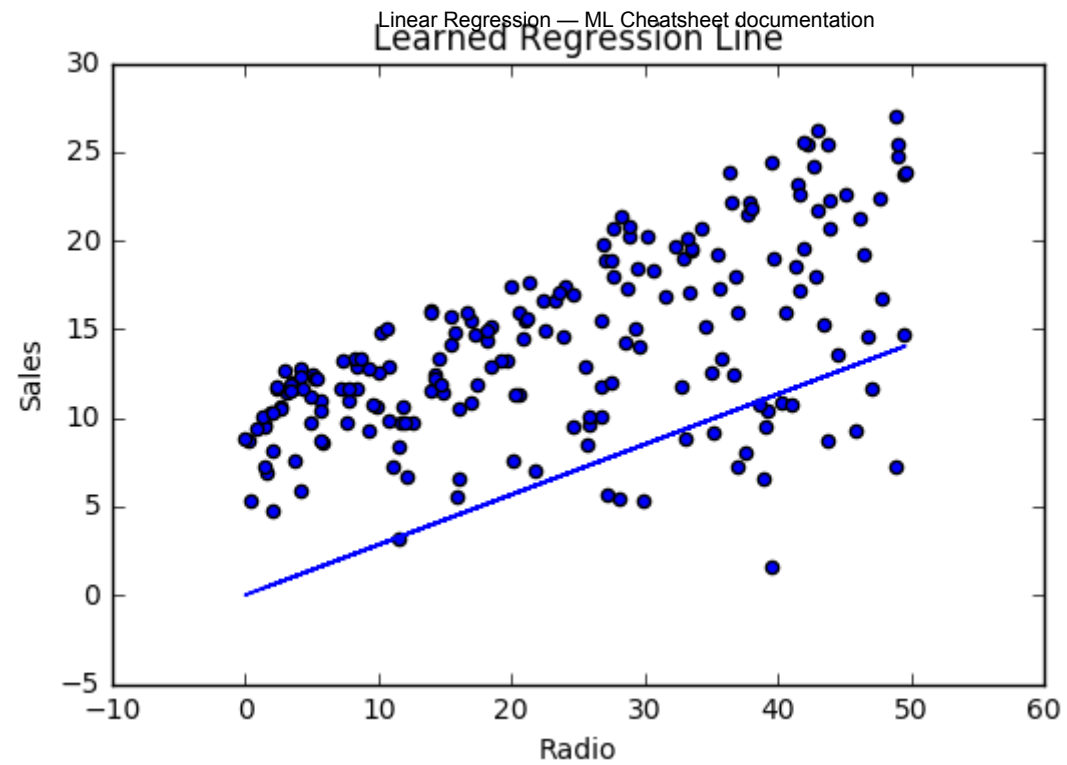
If our model is working, we should see our cost decrease after every iteration.


Logging


iter=1	weight=.03	bias=.0014	cost=197.25
iter=10	weight=.28	bias=.0116	cost=74.65
iter=20	weight=.39	bias=.0177	cost=49.48
iter=30	weight=.44	bias=.0219	cost=44.31
iter=30	weight=.46	bias=.0249	cost=43.28

Visualizing










Cost history



Summary

By learning the best values for weight (.46) and bias (.25), we now have an equation that predicts future sales based on radio advertising investment.

$$Sales = .46Radio + .025$$

How would our model perform in the real world? I'll let you think about it :)

Multivariable regression


Let's say we are given [data](#) on TV, radio, and newspaper advertising spend for a list of companies, and our goal is to predict sales in terms of units sold.

Company	TV	Radio	News	Units
---------	----	-------	------	-------

Amazon	230.1	37.8	69.1	22.1
Google	44.5	39.3	23.1	10.4
Facebook	17.2	45.9	34.7	18.3
Apple	151.5	41.3	13.2	18.5

Growing complexity

As the number of features grows, the complexity of our model increases and it becomes increasingly difficult to visualize, or even comprehend, our data.



One solution is to break the data apart and compare 1-2 features at a time. In this example we explore how Radio and TV investment impacts Sales.

Normalization

As the number of features grows, calculating gradient takes longer to compute. We can speed this up by “normalizing” our input data to ensure all values are within the same range. This is especially important for datasets with high standard deviations or differences in the ranges of the attributes. Our goal now will be to normalize our features so they are all in the range -1 to 1.

Code

```
For each feature column {  
    #1 Subtract the mean of the column (mean normalization)  
    #2 Divide by the range of the column (feature scaling)  
}
```

Our input is a 200 x 3 matrix containing TV, Radio, and Newspaper data. Our output is a normalized matrix of the same shape with all values between -1 and 1.

```
def normalize(features):  
    **  
    features      - (200, 3)  
    features.T    - (3, 200)  
  
    We transpose the input matrix, swapping  
    cols and rows to make vector math easier  
    **  
  
    for feature in features.T:  
        fmean = np.mean(feature)  
        frange = np.amax(feature) - np.amin(feature)  
  
        #Vector Subtraction  
        feature -= fmean  
  
        #Vector Division  
        feature /= frange  
  
    return features
```

Matrix math. Before we continue, it's important to understand basic [Linear Algebra](#) concepts as well as numpy functions like [numpy.dot\(\)](#).

Making predictions

Our predict function outputs an estimate of sales given our current weights (coefficients) and a company's TV, radio, and newspaper spend. Our model will try to identify weight values that most reduce our cost function.

$$Sales = W_1TV + W_2Radio + W_3Newspaper$$

```
def predict(features, weights):  
    **  
    features = (200, 3)  
    weights = (3, 1)  
    predictions = (200,1)  
    **  
    predictions = np.dot(features, weights)  
    return predictions
```

Initialize weights

```
W1 = 0.0  
W2 = 0.0  
W3 = 0.0  
weights = np.array([  
    [W1],  
    [W2],  
    [W3]  
])
```

Cost function

Now we need a cost function to audit how our model is performing. The math is the same, except we swap the $mx + b$ expression for $W_1x_1 + W_2x_2 + W_3x_3$. We also divide the expression by 2 to make derivative calculations simpler.

$$MSE = \frac{1}{2N} \sum_{i=1}^n (y_i - (W_1x_1 + W_2x_2 + W_3x_3))^2$$

```
def cost_function(features, targets, weights):
    """
    features:(200,3)
    targets: (200,1)
    weights:(3,1)
    returns average squared error among predictions
    """
    N = len(targets)

    predictions = predict(features, weights)

    # Matrix math Lets use do this without looping
    sq_error = (predictions - targets)**2

    # Return average squared error among predictions
    return 1.0/(2*N) * sq_error.sum()
```

Gradient descent

Again using the [Chain rule](#) we can compute the gradient—a vector of partial derivatives describing the slope of the cost function for each weight.

$$\begin{aligned} f'(W_1) &= -x_1(y - (W_1x_1 + W_2x_2 + W_3x_3)) \\ f'(W_2) &= -x_2(y - (W_1x_1 + W_2x_2 + W_3x_3)) \\ f'(W_3) &= -x_3(y - (W_1x_1 + W_2x_2 + W_3x_3)) \end{aligned}$$

```

def update_weights(features, targets, weights, lr):
    '''
    Features:(200, 3)
    Targets: (200, 1)
    Weights:(3, 1)
    '''

    predictions = predict(features, weights)

    #Extract our features
    x1 = features[:,0]
    x2 = features[:,1]
    x3 = features[:,2]

    # Use matrix cross product (*) to simultaneously
    # calculate the derivative for each weight
    d_w1 = -x1*(targets - predictions)
    d_w2 = -x2*(targets - predictions)
    d_w3 = -x3*(targets - predictions)

    # Multiply the mean derivative by the Learning rate
    # and subtract from our weights (remember gradient points in direction of steepest ASCENT)
    weights[0][0] -= (lr * np.mean(d_w1))
    weights[1][0] -= (lr * np.mean(d_w2))
    weights[2][0] -= (lr * np.mean(d_w3))

    return weights

```

And that's it! Multivariate linear regression.

Simplifying with matrices

The gradient descent code above has a lot of duplication. Can we improve it somehow? One way to refactor would be to loop through our features and weights—allowing our function to handle any number of features. However there is another even better technique: *vectorized gradient descent*.

Math

We use the same formula as above, but instead of operating on a single feature at a time, we use matrix multiplication to operate on all features and weights simultaneously. We replace the x_i terms with a single feature matrix X .

$$\textit{gradient} = -X(\textit{targets} - \textit{predictions})$$

Code

```

X = [
    [x1, x2, x3]
    [x1, x2, x3]
    .
    .
    .
    [x1, x2, x3]
]

targets = [
    [1],
    [2],
    [3]
]

def update_weights_vectorized(X, targets, weights, lr):
    **
    gradient = X.T * (predictions - targets) / N
    X: (200, 3)
    Targets: (200, 1)
    Weights: (3, 1)
    **

    companies = len(X)

    #1 - Get Predictions
    predictions = predict(X, weights)

    #2 - Calculate error/loss
    error = targets - predictions

    #3 Transpose features from (200, 3) to (3, 200)
    # So we can multiply w the (200,1) error matrix.
    # Returns a (3,1) matrix holding 3 partial derivatives --
    # one for each feature -- representing the aggregate
    # slope of the cost function across all observations
    gradient = np.dot(-X.T, error)

    #4 Take the average error derivative for each feature
    gradient /= companies

    #5 - Multiply the gradient by our Learning rate
    gradient *= lr

    #6 - Subtract from our weights to minimize cost
    weights -= gradient

    return weights

```

Bias term

Our train function is the same as for simple linear regression, however we're going to make one final tweak before running: add a **bias term** to our feature matrix.

In our example, it's very unlikely that sales would be zero if companies stopped advertising. Possible reasons for this might include past advertising, existing customer relationships, retail locations, and salespeople. A bias term will help us capture this base case.

Code

Below we add a constant 1 to our features matrix. By setting this value to 1, it turns our bias term into a constant.


```
bias = np.ones(shape=(len(features),1))
features = np.append(bias, features, axis=1)
```

Model evaluation

After training our model through 1000 iterations with a learning rate of .0005, we finally arrive at a set of weights we can use to make predictions:

$$Sales = 4.7TV + 3.5Radio + .81Newspaper + 13.9$$

Our MSE cost dropped from 110.86 to 6.25.



References

- [1] https://en.wikipedia.org/wiki/Linear_regression
- [2] http://www.holehouse.org/mlclass/04_Linear_Regression_with_multiple_variables.html
- [3] <http://machinelearningmastery.com/simple-linear-regression-tutorial-for-machine-learning>
- [4] <http://people.duke.edu/~rnau/regintro.htm>
- [5] <https://spin.atomicobject.com/2014/06/24/gradient-descent-linear-regression>
- [6] <https://www.analyticsvidhya.com/blog/2015/08/common-machine-learning-algorithms>