

Machine Learning

Machine Learning Resources

[Online Resources \(/machine-learning-resources/\)](/machine-learning-resources/)

[Journal Library \(/machine-learning/journals-library/\)](/machine-learning/journals-library/)

[Datasets for Machine Learning \(/machine-learning/datasets\)](/machine-learning/datasets)

Machine Learning and Econometrics

[Resources \(/machine-learning/econometrics/resources\)](/machine-learning/econometrics/resources)

Supervised Learning Theory

[Overview \(/machine-learning/\)](/machine-learning/)

[One Variable Linear Regression \(/one-variable-linear-regression/\)](/one-variable-linear-regression/)

[Linear Algebra \(/linear-algebra-machine-learning/\)](/linear-algebra-machine-learning/)

[Multiple Variable Linear Regression \(/multi-variable-linear-regression/\)](/multi-variable-linear-regression/)

[Logistic Regression \(/logistic-regression/\)](/logistic-regression/)

[Neural Networks \(Representation\) \(/neural-networks-representation/\)](/neural-networks-representation/)

[Neural Networks \(Learning\) \(/neural-networks-learning/\)](/neural-networks-learning/)

[Applying Machine Learning \(/applying-machine-learning/\)](/applying-machine-learning/)

[Machine Learning Systems Design \(/machine-learning-systems-design/\)](/machine-learning-systems-design/)

[Support Vector Machines \(/machine-learning-svms-support-vector-machines/\)](/machine-learning-svms-support-vector-machines/)

Unsupervised Learning Theory

[Unsupervised Learning \(/machine-learning-unsupervised-learning/\)](/machine-learning-unsupervised-learning/)

[Dimensionality Reduction \(/machine-learning-dimensionality-reduction/\)](/machine-learning-dimensionality-reduction/)

[Anomaly Detection \(/machine-learning-anomaly-detection/\)](/machine-learning-anomaly-detection/)

[Recommender Systems \(/machine-learning-recommender-systems/\)](/machine-learning-recommender-systems/)[Large Scale Machine Learning \(/machine-learning-large-scale/\)](/machine-learning-large-scale/)[Photo OCR \(/machine-learning-photo-ocr/\)](/machine-learning-photo-ocr/)

Reinforcement Learning Theory

[Markov Decision Processes \(/machine-learning-markov-decision-processes/\)](/machine-learning-markov-decision-processes/)[Reinforcement Learning \(/machine-learning-reinforcement-learning/\)](/machine-learning-reinforcement-learning/)[Game Theory \(/machine-learning-game-theory/\)](/machine-learning-game-theory/)

Deep Learning Theory

[Deep Learning Terms \(/machine-learning/deep-learning/terms/\)](/machine-learning/deep-learning/terms/)[Deep Learning Intro \(/machine-learning/deep-learning/intro/\)](/machine-learning/deep-learning/intro/)[Deep Neural Networks Intro \(/machine-learning/deep-learning/neural-nets/\)](/machine-learning/deep-learning/neural-nets/)[Deep Convolutional Networks Intro \(/machine-learning/deep-learning/convnets/\)](/machine-learning/deep-learning/convnets/)

Deep Learning with TensorFlow

[Exploring NotMNIST \(/machine-learning/deep-learning/tensorflow/notmnist/\)](/machine-learning/deep-learning/tensorflow/notmnist/)[Deep Neural Networks \(/machine-learning/deep-learning/tensorflow/deep-neural-nets/\)](/machine-learning/deep-learning/tensorflow/deep-neural-nets/)[Regularization \(/machine-learning/deep-learning/tensorflow/regularization/\)](/machine-learning/deep-learning/tensorflow/regularization/)[Deep Convolutional Networks \(/machine-learning/deep-learning/tensorflow/convnets/\)](/machine-learning/deep-learning/tensorflow/convnets/)

Machine Learning with Scikit-Learn

[Introduction to Machine Learning \(/machine-learning-intro-easy/\)](/machine-learning-intro-easy/)[IPython Introduction \(/ipython-introduction/\)](/ipython-introduction/)[Iris Dataset \(/machine-learning-iris-dataset/\)](/machine-learning-iris-dataset/)[Linear Regression Model \(/machine-learning-linear-regression/\)](/machine-learning-linear-regression/)[Linear Regression Model Evaluation \(/machine-learning-evaluate-linear-regression-model/\)](/machine-learning-evaluate-linear-regression-model/)[Polynomial Regression \(/machine-learning-polynomial-regression/\)](/machine-learning-polynomial-regression/)[Vectorization, Multinomial Naive Bayes Classifier and Evaluation \(/machine-learning-multinomial-naive-bayes-vectorization/\)](/machine-learning-multinomial-naive-bayes-vectorization/)

[Gaussian Naive Bayes \(/machine-learning-gaussian-naive-bayes/\)](#)[K-nearest Neighbors \(KNN\) Classification Model \(/machine-learning-k-nearest-neighbors-knn/\)](#)[Ensemble Learning and Adaboost \(/machine-learning-ensemble-of-learners-adaboost/\)](#)[Decision Trees \(/machine-learning-decision-trees/\)](#)[Support Vector Machines \(/machine-learning-svms/\)](#)[Clustering with KMeans \(/machine-learning-clustering-kmeans/\)](#)[Dimensionality Reduction and Feature Transformation \(/machine-learning-dimensionality-reduction-feature-transform/\)](#)[Feature Engineering and Scaling \(/machine-learning-feature_engineering_scaling/\)](#)[Cross-Validation for Parameter Tuning, Model Selection, and Feature Selection \(/machine-learning-cross-validation/\)](#)[Efficiently Searching Optimal Tuning Parameters \(/machine-learning-efficiently-search-tuning-param/\)](#)[Evaluating a Classification Model \(/machine-learning-evaluate-classification-model/\)](#)[One Hot Encoding \(/machinelearning-one-hot-encoding/\)](#)[F1 Score \(/machinelearning-f1-score/\)](#)[Learning Curve \(/machinelearning-learning-curve/\)](#)

Machine Learning Projects

[Titanic Survival Data Exploration \(/machine-learning-project-titanic-survival/\)](#)[Boston House Prices Prediction and Evaluation \(Model Evaluation and Prediction\) \(/machine-learning-project-boston-home-prices/\)](#)[Building a Student Intervention System \(Supervised Learning\) \(/machine-learning-project-student-intervention/\)](#)[Identifying Customer Segments \(Unsupervised Learning\) \(/machine-learning-project-customer-segments/\)](#)[Training a Smart Cab \(Reinforcement Learning\) \(/machine-learning-proj-smart-cab/\)](#)

Cross-Validation

Summary: Cross-Validation for Parameter Tuning, Model Selection, and Feature Selection

Topics

1. Review of model evaluation procedures
2. Steps for K-fold cross-validation
3. Comparing cross-validation to train/test split
4. Cross-validation recommendations
5. Cross-validation example: parameter tuning
6. Cross-validation example: model selection
7. Cross-validation example: feature selection
8. Improvements to cross-validation
9. Resources

This tutorial is derived from Data School's Machine Learning with scikit-learn tutorial. I added my own notes so anyone, including myself, can refer to this tutorial without watching the videos.

1. Review of model evaluation procedures

Motivation: Need a way to choose between machine learning models

- Goal is to estimate likely performance of a model on **out-of-sample data**

Initial idea: Train and test on the same data

- But, maximizing **training accuracy** rewards overly complex models which **overfit** the training data

Alternative idea: Train/test split

- Split the dataset into two pieces, so that the model can be trained and tested on **different data**
- **Testing accuracy** is a better estimate than training accuracy of out-of-sample performance
- Problem with train/test split
 - It provides a **high variance** estimate since changing which observations happen to be in the testing set can significantly change testing accuracy
 - Testing accuracy can change a lot depending on a which observation happen to be in the testing set

```
In [1]: from sklearn.datasets import load_iris
        from sklearn.cross_validation import train_test_split
        from sklearn.neighbors import KNeighborsClassifier
        from sklearn import metrics
```

```
In [2]: # read in the iris data
        iris = load_iris()

        # create X (features) and y (response)
        X = iris.data
        y = iris.target
```

```
In [3]: # use train/test split with different random_state values  
# we can change the random_state values that changes the accuracy scores  
# the accuracy changes a lot  
# this is why testing accuracy is a high-variance estimate  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=6)  
  
# check classification accuracy of KNN with K=5  
knn = KNeighborsClassifier(n_neighbors=5)  
knn.fit(X_train, y_train)  
y_pred = knn.predict(X_test)  
metrics.accuracy_score(y_test, y_pred)
```

```
Out[3]: 0.97368421052631582
```

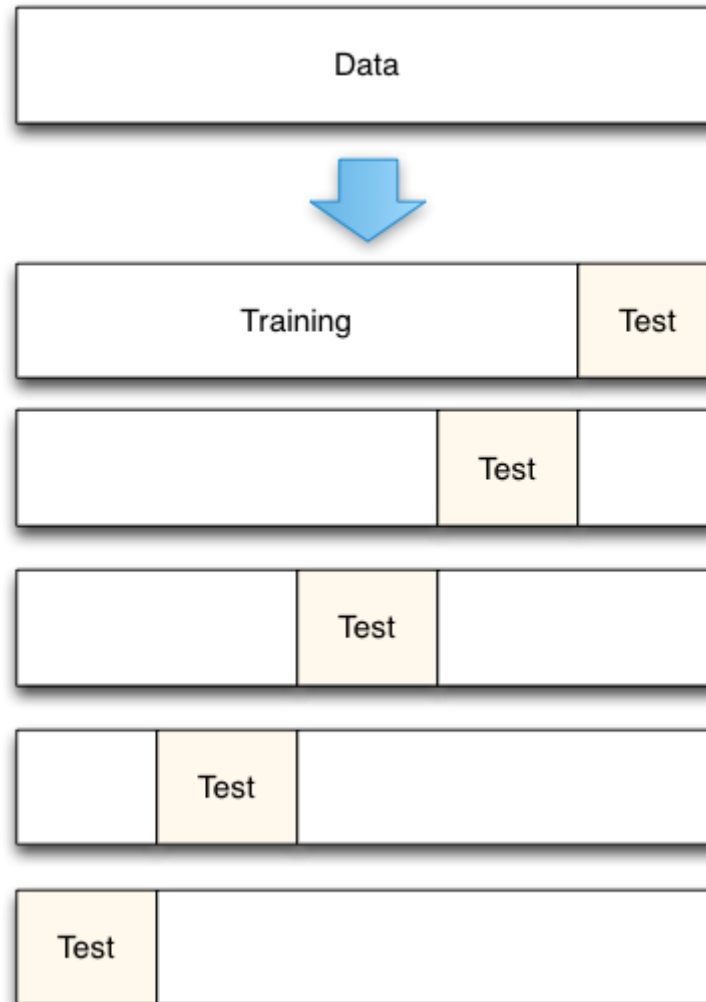
Question: What if we created a bunch of train/test splits, calculated the testing accuracy for each, and averaged the results together?

Answer: That's the essence of cross-validation!

2. Steps for K-fold cross-validation

1. Split the dataset into K **equal** partitions (or "folds")
 - So if $k = 5$ and dataset has 150 observations
 - Each of the 5 folds would have 30 observations
2. Use fold 1 as the **testing set** and the union of the other folds as the **training set**
 - Testing set = 30 observations (fold 1)
 - Training set = 120 observations (folds 2-5)
3. Calculate **testing accuracy**
4. Repeat steps 2 and 3 K times, using a **different fold** as the testing set each time
 - We will repeat the process 5 times
 - 2nd iteration
 - fold 2 would be the testing set
 - union of fold 1, 3, 4, and 5 would be the training set
 - 3rd iteration
 - fold 3 would be the testing set
 - union of fold 1, 2, 4, and 5 would be the training set
 - And so on...
5. Use the **average testing accuracy** as the estimate of out-of-sample accuracy

Diagram of **5-fold cross-validation**:




```
In [4]: # simulate splitting a dataset of 25 observations into 5 folds
from sklearn.cross_validation import KFold
kf = KFold(25, n_folds=5, shuffle=False)

# print the contents of each training and testing set
# ^ - forces the field to be centered within the available space
# .format() - formats the string similar to %s or %n
# enumerate(sequence, start=0) - returns an enumerate object
print('{} {:^61} {}'.format('Iteration', 'Training set observations', 'Testing set observations'))
for iteration, data in enumerate(kf, start=1):
    print('{!s:^9} {} {!s:^25}'.format(iteration, data[0], data[1]))
```

Iteration	Training set observations	Testing set observations
1	[5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24]	[0 1 2 3 4]
2	[0 1 2 3 4 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24]	[5 6 7 8 9]
3	[0 1 2 3 4 5 6 7 8 9 15 16 17 18 19 20 21 22 23 24]	[10 11 12 13 14]
4	[0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 20 21 22 23 24]	[15 16 17 18 19]
5	[0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19]	[20 21 22 23 24]

- Dataset contains **25 observations** (numbered 0 through 24)
- 5-fold cross-validation, thus it runs for **5 iterations**
- For each iteration, every observation is either in the training set or the testing set, **but not both**
- Every observation is in the testing set **exactly once**

3. Comparing cross-validation to train/test split

Advantages of cross-validation:

- More accurate estimate of out-of-sample accuracy
- More "efficient" use of data
 - This is because every observation is used for both training and testing

Advantages of train/test split:

- Runs K times faster than K-fold cross-validation
 - This is because K-fold cross-validation repeats the train/test split K-times
- Simpler to examine the detailed results of the testing process

4. Cross-validation recommendations

1. K can be any number, but **K=10** is generally recommended
 - This has been shown experimentally to produce the best out-of-sample estimate
2. For classification problems, **stratified sampling** is recommended for creating the folds
 - Each response class should be represented with equal proportions in each of the K folds
 - If dataset has 2 response classes
 - Spam/Ham
 - 20% observation = ham
 - Each cross-validation fold should consist of exactly 20% ham
 - scikit-learn's `cross_val_score` function does this by default

5. Cross-validation example: parameter tuning

Goal: Select the best tuning parameters (aka "hyperparameters") for KNN on the iris dataset

- We want to choose the best tuning parameters that best generalize the data

```
In [5]: from sklearn.cross_validation import cross_val_score
```

```
In [6]: # 10-fold cross-validation with K=5 for KNN (the n_neighbors parameter)
# k = 5 for KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=5)

# Use cross_val_score function
# We are passing the entirety of X and y, not X_train or y_train, it takes care of splitting the dat
# cv=10 for 10 folds
# scoring='accuracy' for evaluation metric - although they are many
scores = cross_val_score(knn, X, y, cv=10, scoring='accuracy')
print(scores)

[ 1.          0.93333333  1.          1.          0.86666667  0.93333333
  0.93333333  1.          1.          1.          ]
```

- In the first iteration, the accuracy is 100%
- Second iteration, the accuracy is 93% and so on

cross_val_score executes the first 4 steps of k-fold cross-validation steps which I have broken down to 7 steps here in detail

1. Split the dataset (X and y) into K=10 **equal** partitions (or "folds")
2. Train the KNN model on union of folds 2 to 10 (training set)
3. Test the model on fold 1 (testing set) and calculate testing accuracy
4. Train the KNN model on union of fold 1 and fold 3 to 10 (training set)
5. Test the model on fold 2 (testing set) and calculate testing accuracy
6. It will do this on 8 more times
7. When finished, it will return the 10 testing accuracy scores as a numpy array

```
In [7]: # use average accuracy as an estimate of out-of-sample accuracy
# numpy array has a method mean()
print(scores.mean())
```

0.966666666667

Our goal here is to find the optimal value of K

```
In [8]: # search for an optimal value of K for KNN

# range of k we want to try
k_range = range(1, 31)
# empty list to store scores
k_scores = []

# 1. we will loop through reasonable values of k
for k in k_range:
    # 2. run KNeighborsClassifier with k neighbours
    knn = KNeighborsClassifier(n_neighbors=k)
    # 3. obtain cross_val_score for KNeighborsClassifier with k neighbours
    scores = cross_val_score(knn, X, y, cv=10, scoring='accuracy')
    # 4. append mean of scores for k neighbors to k_scores list
    k_scores.append(scores.mean())

print(k_scores)
```

```
[0.95999999999999996, 0.95333333333333337, 0.96666666666666656, 0.96666666666666656, 0.96666666666666679, 0.96666666666666679, 0.96666666666666679, 0.96666666666666679, 0.97333333333333338, 0.96666666666666679, 0.96666666666666666, 0.97333333333333338, 0.98000000000000009, 0.97333333333333338, 0.97333333333333338, 0.97333333333333338, 0.97333333333333338, 0.98000000000000009, 0.97333333333333338, 0.98000000000000009, 0.96666666666666656, 0.96666666666666656, 0.97333333333333338, 0.95999999999999996, 0.96666666666666656, 0.95999999999999996, 0.96666666666666656, 0.95333333333333337, 0.95333333333333337, 0.95333333333333337]
```

```
In [9]: # in essence, this is basically running the k-fold cross-validation method 30 times because we want to run through  
K values from 1 to 30  
# we should have 30 scores here  
print('Length of list', len(k_scores))  
print('Max of list', max(k_scores))
```


Length of list 30

Max of list 0.98

```
In [10]: # plot how accuracy changes as we vary k
import matplotlib.pyplot as plt
%matplotlib inline

# plot the value of K for KNN (x-axis) versus the cross-validated accuracy (y-axis)
# plt.plot(x_axis, y_axis)
plt.plot(k_range, k_scores)
plt.xlabel('Value of K for KNN')
plt.ylabel('Cross-validated accuracy')
```

Out[10]: <matplotlib.text.Text at 0x111cb07b8>



The maximum cv accuracy occurs from $k=13$ to $k=20$

- The general shape of the curve is an upside down yield
 - This is quite typical when examining the model complexity and accuracy
 - This is an example of bias-variance trade off
 - Low values of k (low bias, high variance)
 - The 1-Nearest Neighbor classifier is the most complex nearest neighbor model
 - It has the most jagged decision boundary, and is most likely to overfit
 - High values of k (high bias, low variance)
 - underfit
 - Best value is the middle of k (most likely to generalize out-of-sample data)
 - just right
- The best value of k
 - Higher values of k produce less complex model
 - So we will choose 20 as our best KNN model

6. Cross-validation example: model selection

Goal: Compare the best KNN model with logistic regression on the iris dataset


```
In [11]: # 10-fold cross-validation with the best KNN model
knn = KNeighborsClassifier(n_neighbors=20)

# Instead of saving 10 scores in object named score and calculating mean
# We're just calculating the mean directly on the results
print(cross_val_score(knn, X, y, cv=10, scoring='accuracy').mean())

0.98
```

```
In [12]: # 10-fold cross-validation with logistic regression
from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression()
print(cross_val_score(logreg, X, y, cv=10, scoring='accuracy').mean())

0.953333333333
```

We can conclude that KNN is likely a better choice than logistic regression

7. Cross-validation example: feature selection

Goal: Select whether the Newspaper feature should be included in the linear regression model on the advertising dataset

```
In [13]: import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
```

```
In [14]: # read in the advertising dataset
data = pd.read_csv('http://www-bcf.usc.edu/~gareth/ISL/Advertising.csv (http://www-bcf.usc.edu/~gareth/ISL/Advertising.csv)', index_col=0)
```

```
In [15]: # create a Python List of three feature names
feature_cols = ['TV', 'Radio', 'Newspaper']

# use the list to select a subset of the DataFrame (X)
X = data[feature_cols]

# select the Sales column as the response (y)
# since we're selecting only one column, we can select the attribute using .attribute
y = data.Sales
```

```
In [16]: # 10-fold cross-validation with all three features
# instantiate model
lm = LinearRegression()

# store scores in scores object
# we can't use accuracy as our evaluation metric since that's only relevant for classification problems
# RMSE is not directly available so we will use MSE
scores = cross_val_score(lm, X, y, cv=10, scoring='mean_squared_error')
print(scores)

[-3.56038438 -3.29767522 -2.08943356 -2.82474283 -1.3027754  -1.74163618
 -8.17338214 -2.11409746 -3.04273109 -2.45281793]
```

MSE should be positive

- But why is the MSE here negative?
- MSE is a loss function
 - It is something we want to minimize
 - A design decision was made so that the results are made negative
 - The best results would be the largest number (the least negative) so we can still maximize similar to classification accuracy
- Classification Accuracy is a reward function
 - It is something we want to maximize

```
In [17]: # fix the sign of MSE scores  
mse_scores = -scores  
print(mse_scores)
```

```
[ 3.56038438  3.29767522  2.08943356  2.82474283  1.3027754   1.74163618  
 8.17338214  2.11409746  3.04273109  2.45281793]
```

```
In [18]: # convert from MSE to RMSE  
rmse_scores = np.sqrt(mse_scores)  
print(rmse_scores)
```

```
[ 1.88689808  1.81595022  1.44548731  1.68069713  1.14139187  1.31971064  
 2.85891276  1.45399362  1.7443426   1.56614748]
```

```
In [19]: # calculate the average RMSE  
print(rmse_scores.mean())
```

```
1.69135317081
```

```
In [20]: # 10-fold cross-validation with two features (excluding Newspaper)
feature_cols = ['TV', 'Radio']
X = data[feature_cols]
print(np.sqrt(-cross_val_score(lm, X, y, cv=10, scoring='mean_squared_error')).mean())
```

1.67967484191

Without Newspaper

- Average RMSE = 1.68
- lower number than with model with Newspaper
 - RMSE is something we want to minimize
 - So the model excluding Newspaper is a better model

8. Improvements to cross-validation

Repeated cross-validation

- Repeat cross-validation multiple times (with **different random splits** of the data) and average the results
- More reliable estimate of out-of-sample performance by **reducing the variance** associated with a single trial of cross-validation

Creating a hold-out set

- "Hold out" a portion of the data **before** beginning the model building process
- Locate the best model using cross-validation on the remaining data, and test it **using the hold-out set**
- More reliable estimate of out-of-sample performance since hold-out set is **truly out-of-sample**

Feature engineering and selection within cross-validation iterations

- Normally, feature engineering and selection occurs **before** cross-validation
- Instead, perform all feature engineering and selection **within each cross-validation iteration**
- More reliable estimate of out-of-sample performance since it **better mimics** the application of the model to out-of-sample data

9. Resources

- scikit-learn documentation: [Cross-validation](http://scikit-learn.org/stable/modules/cross_validation.html) (http://scikit-learn.org/stable/modules/cross_validation.html), [Model evaluation](http://scikit-learn.org/stable/modules/model_evaluation.html) (http://scikit-learn.org/stable/modules/model_evaluation.html)
- scikit-learn issue on GitHub: [MSE is negative when returned by cross_val_score](https://github.com/scikit-learn/scikit-learn/issues/2439) (<https://github.com/scikit-learn/scikit-learn/issues/2439>)
- Section 5.1 of [An Introduction to Statistical Learning](http://www-bcf.usc.edu/~gareth/ISL/) (<http://www-bcf.usc.edu/~gareth/ISL/>) (11 pages) and related videos: [K-fold and leave-one-out cross-validation](https://www.youtube.com/watch?v=nZAM5OXrktY) (<https://www.youtube.com/watch?v=nZAM5OXrktY>) (14 minutes), [Cross-validation the right and wrong ways](https://www.youtube.com/watch?v=S06JpVoNaA0) (<https://www.youtube.com/watch?v=S06JpVoNaA0>) (10 minutes)
- Scott Fortmann-Roe: [Accurately Measuring Model Prediction Error](http://scott.fortmann-roe.com/docs/MeasuringError.html) (<http://scott.fortmann-roe.com/docs/MeasuringError.html>)
- Machine Learning Mastery: [An Introduction to Feature Selection](http://machinelearningmastery.com/an-introduction-to-feature-selection/) (<http://machinelearningmastery.com/an-introduction-to-feature-selection/>)
- Harvard CS109: [Cross-Validation: The Right and Wrong Way](https://github.com/cs109/content/blob/master/lec_10_cross_val.ipynb) (https://github.com/cs109/content/blob/master/lec_10_cross_val.ipynb)
- Journal of Cheminformatics: [Cross-validation pitfalls when selecting and assessing regression and classification models](http://www.jcheminf.com/content/pdf/1758-2946-6-10.pdf) (<http://www.jcheminf.com/content/pdf/1758-2946-6-10.pdf>)

Tags:

machine_learning (/tag_machine_learning)

1 Comment Ritchie Ng | Deep Learning & Computer Vision Engineer

1 Login ▾

♥ Recommend 4  Tweet  Share

Sort by Best ▾



Join the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS ?

**KELVIN TAN** • a year ago

I am using Python 3.6, and having this error:

C:\Users\xxx\Anaconda3\lib\site-packages\sklearn\metrics\scorer.py:100: DeprecationWarning: Scoring method mean_squared_error was renamed to neg_mean_squared_error in version 0.18 and will be removed in 0.20.

sample_weight=sample_weight)

when run the line:

scores = cross_val_score(lm, X, y, cv=10, scoring='mean_squared_error'), so I change to

scores = cross_val_score(lm, X, y, cv=10, scoring='neg_mean_squared_error'), and it works.

also 'Radio' should be 'radio' and 'Newspaper' should be 'newspaper'.

y = data.Sales should be y = data.sales

Can you also explain why 31 in

k_range = range(1, 31) ?

^ | v • Reply • Share ›

ALSO ON RITCHIE NG | DEEP LEARNING & COMPUTER VISION ENGINEER

Deep Neural Networks with TensorFlow | Machine Learning, Deep Learning, Reinforcement Learning ...

9 comments • 3 years ago



Kavish Rastogi — Dear Ritchie and all the techies here, I am new to deep learning. I have a doubt here .When we are applying S.G.D. with num_steps=3001 , How many ...

Machine Learning Photo OCR | Machine Learning, Deep Learning, and Computer Vision

I am an NVIDIA Deep Learning Institute Instructor | Machine Learning, Deep Learning, and Computer ...

1 comment • 2 years ago



Aitor Lopez Beltran — First of all, congratulations! Excuse me, but I'm having a hard time finding out how to get certified as a DLI Instructor. Could you help me, please? ...

Exploring Series | Machine Learning, Deep Learning, Reinforcement Learning and Artificial Intelligence ...

4 comments • 3 years ago






Aayush Agarwal — How to implement in code? Any source code available?

1 comment • 3 years ago



Kurt Hjortdahl — If you need to group by intervals, check out ...

 **Subscribe**  **Add Disqus to your site**[Add DisqusAdd](#)  **Disqus' Privacy Policy**[Privacy Policy](#)[Privacy](#)[Privacy](#)

© 2019 Ritchie Ng. All rights reserved.

Site last updated: May 22, 2019

[Github \(https://github.com/ritchieng\)](https://github.com/ritchieng) | [Linkedin \(https://www.linkedin.com/in/ritchieng\)](https://www.linkedin.com/in/ritchieng) | [Facebook \(https://www.facebook.com/ritchiengz\)](https://www.facebook.com/ritchiengz) | [Twitter \(https://twitter.com/ritchieng\)](https://twitter.com/ritchieng) | [Tech in Asia \(https://www.techinasia.com/profile/ritchieng\)](https://www.techinasia.com/profile/ritchieng)