



[Home](#) [Installation](#)
[Documentation](#)
[Examples](#)

sklearn.linear_model.LinearRegression

```
class sklearn.linear_model.LinearRegression(fit_intercept=True, normalize=False, copy_X=True, n_jobs=None)
```

[\[source\]](#)

»

Ordinary least squares Linear Regression.

Parameters: **fit_intercept** : *boolean, optional, default True*

whether to calculate the intercept for this model. If set to False, no intercept will be used in calculations (e.g. data is expected to be already centered).

normalize : *boolean, optional, default False*

This parameter is ignored when `fit_intercept` is set to False. If True, the regressors X will be normalized before regression by subtracting the mean and dividing by the l2-norm. If you wish to standardize, please use [sklearn.preprocessing.StandardScaler](#) before calling `fit` on an estimator with `normalize=False`.

copy_X : *boolean, optional, default True*

If True, X will be copied; else, it may be overwritten.

n_jobs : *int or None, optional (default=None)*

The number of jobs to use for the computation. This will only provide speedup for `n_targets > 1` and sufficient large problems. None means 1 unless in a [joblib.parallel_backend](#) context. -1 means using all processors. See [Glossary](#) for more details.

Attributes: **coef_** : *array, shape (n_features,) or (n_targets, n_features)*

Estimated coefficients for the linear regression problem. If multiple targets are passed during the fit (y 2D), this is a 2D array of shape (n_targets, n_features), while if only one target is passed, this is a 1D array of length n_features.

intercept_ : *array*

Independent term in the linear model.

Notes

From the implementation point of view, this is just plain Ordinary Least Squares (`scipy.linalg.lstsq`) wrapped as a predictor object.

Examples

```
>>> import numpy as np
>>> from sklearn.linear_model import LinearRegression
>>> X = np.array([[1, 1], [1, 2], [2, 2], [2, 3]])
>>> # y = 1 * x_0 + 2 * x_1 + 3
>>> y = np.dot(X, np.array([1, 2])) + 3
>>> reg = LinearRegression().fit(X, y)
>>> reg.score(X, y)
1.0
>>> reg.coef_
array([1., 2.])
>>> reg.intercept_
3.0000...
>>> reg.predict(np.array([[3, 5]]))
array([16.])
```

Methods

<code>fit(self, X, y[, sample_weight])</code>	Fit linear model.
<code>get_params(self[, deep])</code>	Get parameters for this estimator.
<code>predict(self, X)</code>	Predict using the linear model
<code>score(self, X, y[, sample_weight])</code>	Returns the coefficient of determination R^2 of the prediction.
<code>set_params(self, **params)</code>	Set the parameters of this estimator.

`__init__(self, fit_intercept=True, normalize=False, copy_X=True, n_jobs=None)`

[\[source\]](#)

`fit(self, X, y, sample_weight=None)`

[\[source\]](#)

Fit linear model.

Parameters: **X** : array-like or sparse matrix, shape (n_samples, n_features)

Training data

y : array_like, shape (n_samples, n_targets)

Target values. Will be cast to X's dtype if necessary

sample_weight : *numpy array of shape [n_samples]*

Individual weights for each sample

New in version 0.17: parameter *sample_weight* support to LinearRegression.

Returns: **self** : *returns an instance of self.*

»

get_params(*self*, *deep*=True)

[\[source\]](#)

Get parameters for this estimator.

Parameters: **deep** : *boolean, optional*

If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns: **params** : *mapping of string to any*

Parameter names mapped to their values.

predict(*self*, *X*)

[\[source\]](#)

Predict using the linear model

Parameters: **X** : *array_like or sparse matrix, shape (n_samples, n_features)*

Samples.

Returns: **C** : *array, shape (n_samples,)*

Returns predicted values.

score(*self*, *X*, *y*, *sample_weight*=None)

[\[source\]](#)

Returns the coefficient of determination R^2 of the prediction.

The coefficient R^2 is defined as $(1 - u/v)$, where u is the residual sum of squares $((y_{\text{true}} - y_{\text{pred}})^2).sum()$ and v is the total sum of squares $((y_{\text{true}} - y_{\text{true}.mean()})^2).sum()$. The best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of y , disregarding the input features, would get a R^2 score of 0.0.

Parameters: **X** : array-like, shape = (n_samples, n_features)

Test samples. For some estimators this may be a precomputed kernel matrix instead, shape = (n_samples, n_samples_fitted], where n_samples_fitted is the number of samples used in the fitting for the estimator.

y : array-like, shape = (n_samples) or (n_samples, n_outputs)

True values for X.

sample_weight : array-like, shape = [n_samples], optional

Sample weights.

Returns: **score** : float

R^2 of self.predict(X) wrt. y.

Notes

The R^2 score used when calling `score` on a regressor will use `multioutput='uniform_average'` from version 0.23 to keep consistent with `metrics.r2_score`. This will influence the `score` method of all the multioutput regressors (except for `multioutput.MultiOutputRegressor`). To specify the default value manually and avoid the warning, please either call `metrics.r2_score` directly or make a custom scorer with `metrics.make_scorer` (the built-in scorer 'r2' uses `multioutput='uniform_average'`).

```
set_params(self, **params)
```




[\[source\]](#)

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as pipelines). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Returns: self

Examples using `sklearn.linear_model.LinearRegression`







»

Isotonic Regression

Face completion with a multi-output estimators

Plot individual and voting regression predictions

Ordinary Least Squares and Ridge Regression Variance







Logistic function

Linear Regression Example

Robust linear model estimation using RANSAC

Sparsity Example: Fitting only features 1 and 2




Theil-Sen Regression


Robust linear estimator fitting

Automatic Relevance Determination Regression (ARD)


Bayesian Ridge Regression



» Plotting Cross-Validated Predictions



Underfitting vs. Overfitting



Using KBinsDiscretizer to discretize continuous features